

COMP 540 Final Note

1 Linear models for regression

1.1 Least squares regression

- Loss function: $J(\theta) = \frac{1}{2m} \sum_{i=1}^m (y^{(i)} - h_{\theta}(x^{(i)}))^2$
Vectorized form: $J(\theta) = \frac{1}{2m} (X\theta - y)^T (X\theta - y)$
- Gradient: $\nabla_{\theta} J(\theta) = \frac{1}{m} X^T (X\theta - y)$
- Closed form solution: $\hat{\theta} = (X^T X)^{-1} X^T y$
- Assume $y^{(i)} = (\theta^*)^T x^{(i)} + \epsilon$, where $E[\epsilon] = 0, Var[\epsilon] = \sigma^2$:
 $E[\hat{\theta}] = \theta^*, Var[\hat{\theta}] = (X^T X)^{-1} \sigma^2$.
- \equiv MLE on θ with normal distributed error ϵ .

1.2 L2-regularization: ridge regression

- Loss function:
 $J(\theta) = \frac{1}{2m} \sum_{i=1}^m (y^{(i)} - h_{\theta}(x^{(i)}))^2 + \frac{\lambda}{2m} \sum_{j=1}^d \theta_j^2$
- Closed form solution: $\hat{\theta} = (X^T X + \lambda I)^{-1} X^T y$
- \equiv MAP on θ with prior: $\theta \sim \mathcal{N}(0, \alpha^2 I)$, with $\frac{\lambda}{m} = \frac{\sigma^2}{\alpha^2}$.

1.3 L1-regularization: lasso regression

- Loss function:
 $J(\theta) = \frac{1}{2m} \sum_{i=1}^m (y^{(i)} - h_{\theta}(x^{(i)}))^2 + \frac{\lambda}{2m} \sum_{j=1}^d |\theta_j|$
- \equiv MAP on θ with prior: $\theta_j \sim \text{Laplace}(0, \alpha), \forall j$

1.4 Locally weighted linear regression

- Loss function: $J(\theta) = \frac{1}{2m} \sum_{i=0}^m w^{(i)} (y^{(i)} - \theta^T x^{(i)})^2$
where $w^{(i)} = \exp\left(-\frac{(x-x^{(i)})^T (x-x^{(i)})}{2\sigma^2}\right)$
Vectorized form: $J(\theta) = \frac{1}{2m} (X\theta - y)^T W (X\theta - y)$
- Gradient: $\nabla_{\theta} J(\theta) = \frac{1}{m} X^T W (X\theta - y)$
- Closed form solution: $\hat{\theta} = (X^T W X)^{-1} X^T W y$
- Non-parametric method.

2 Linear models for classification

2.1 Discriminative models for classification

- $P(y=1|x) = h_{\theta}(x) = \frac{1}{1+\exp(-\theta^T x)}$
 $\Rightarrow \log \frac{P(y=1|x)}{P(y=0|x)} = \theta^T x$
- Loss (cross-entropy) function: (*convex & has a global minimum*) $J(\theta) = -\frac{1}{m} \sum_{i=0}^m y^{(i)} \log h_{\theta}(x^{(i)}) + (1 - y^{(i)}) \log(1 - h_{\theta}(x^{(i)}))$
- L2-regularization: $J_{reg}(\theta) = J(\theta) + \frac{\lambda}{m} \sum_{j=1}^d \theta_j^2$
- L1-regularization: $J_{reg}(\theta) = J(\theta) + \frac{\lambda}{m} \sum_{j=1}^d |\theta_j|$

2.2 Generative models for classification

2.2.1 Gaussian discriminant analysis (GDA)

- Assumptions: $y \sim \text{Bernoulli}(\phi)$
 $x|y=0 \sim \mathcal{N}(\mu_0, \Sigma)$ & $x|y=1 \sim \mathcal{N}(\mu_1, \Sigma)$
- Likelihood: $\mathcal{L}(D) = \prod_{i=1}^m P(x^{(i)}, y^{(i)}; \phi, \mu_0, \mu_1, \Sigma)$
 $= \prod_{i=1}^m \phi^{y^{(i)}} (1 - \phi)^{(1-y^{(i)})} \mathcal{N}(x^{(i)}|\mu_1, \Sigma)^{y^{(i)}} \mathcal{N}(x^{(i)}|\mu_0, \Sigma)^{(1-y^{(i)})}$
- Estimation for parameters:
 $\phi = \frac{1}{m} \sum_{i=1}^m y^{(i)}$
 $\mu_1 = \frac{\sum_{i=1}^m y^{(i)} x^{(i)}}{\sum_{i=1}^m y^{(i)}}, \mu_0 = \frac{\sum_{i=1}^m (1-y^{(i)}) x^{(i)}}{\sum_{i=1}^m (1-y^{(i)})}$
 $\Sigma = \frac{1}{m} \sum_{i=1}^m (x - \mu_{y^{(i)}})(x - \mu_{y^{(i)}})^T$
- Linear decision boundaries when same Σ ; quadratic boundaries when each class has its own Σ .

2.2.2 Naive Bayes models

- Assumptions: $P(x|y) = \prod_{j=1}^d P(x_j|y)$
- Bernoulli Naive Bayes models are estimated using counts, regularize using Beta prior (\equiv pre-count).

2.3 Model criteria

- False negative: positive predicted to be negative.
- False positive: negative predicted to be positive.
- specificity = $P(y_{pred} = 0|y = 0) = \frac{TN}{FP+TN}$
- sensitivity = $P(y_{pred} = 1|y = 1) = \frac{TP}{FN+TP}$
- true positive rate (TPR) = sensitivity
- false positive rate (FPR) = 1 - specificity
- ROC curve represents FPR and TPR as a function of classification threshold. $0.5 \leq (\text{area under curve}) \leq 1.0$.

2.4 Multiclass classification

2.4.1 One vs. All (OVA) & One vs. One (OVO)

- OVA: not theoretically justified; simple and widely used.
- OVO: needs $O(K^2)$ classifiers for K classes; overfitting!

2.4.2 Softmax

- Log-likelihood:
 $\ell(\mathcal{D}) = \frac{1}{m} \sum_{i=1}^m \sum_{c=1}^K I(y^{(i)} = c) \log \frac{\exp(\theta^{(c)T} x^{(i)})}{\sum_{c'} \exp(\theta^{(c')T} x^{(i)})}$
- Regularized loss: $J(\theta) = -\ell(\mathcal{D}) + \frac{\lambda}{2m} \sum_{j=1}^d \sum_{c=1}^K \theta_j^{(c)2}$
- Gradient: $\nabla_{\theta} J(\theta) = -\frac{1}{m} \sum_{i=1}^m x^{(i)} (I\{y^{(i)} = c\} - P(y^{(i)} = c|x^{(i)}; \theta)) + \frac{\lambda}{m} \sum_{j=1}^d \theta_j^{(c)}$

3 Kernel methods

3.1 Kernel functions

- Gaussian kernel: $\kappa(x, x') = \exp\left(-\frac{\|x-x'\|^2}{2\sigma^2}\right)$
- Polynomial kernel: $\kappa(x, x') = (1 + x^T x')^p$
- Mercer's theorem (\Leftrightarrow valid kernel): Gram matrix K whose elements are $\kappa(x^{(i)}, x^{(j)})$, $1 \leq i, j \leq m$, should be positive definite for all possible $\{x^{(i)} | 1 \leq i \leq m\}$
 $\Leftrightarrow \exists \phi$ s.t. $\kappa(x, x') = \phi(x)^T \phi(x')$

3.2 Perceptron

- Prediction: $h_{\theta}(x) = \text{sign}(\theta^T x)$
- Update rule: $\theta \leftarrow \theta + \eta x^{(i)} y^{(i)}$, when $h_{\theta}(x^{(i)}) y^{(i)} = -1$
- Convergence bounds: Let $\|x^{(i)}\| \leq R, \forall 1 \leq i \leq m$, the perceptron converges in at most $\frac{R^2 \|\theta^*\|^2}{\gamma^2}$ updates, where $\gamma > 0, y^{(i)}(\theta^T x^{(i)}) \geq \gamma, \forall 1 \leq i \leq m$. (*Appendix*)
- Kernalized version: (if $\eta = 1, \theta = \sum_{(x,y) \in \mathcal{D}_{mistake}} xy$)
 $\hat{y} = \text{sign}(\sum_{(x^{(i)}, y^{(i)}) \in \mathcal{D}} \alpha^{(i)} \langle x^{(i)}, x \rangle)$, in training:
update $\alpha^{(i)} \leftarrow \alpha^{(i)} + y^{(i)}$, when $y\hat{y} = -1$

3.3 Support vector machine (SVM)

3.3.1 Maximize margin

- Optimization problem: $\min_{\theta, \theta_0} \frac{1}{2} \|\theta\|^2$, subject to $y^{(i)}(\theta^T x^{(i)} + \theta_0) \geq 1, \forall 1 \leq i \leq m$
- Dual problem:
 $\max_{\alpha} \sum_{i=1}^m \alpha_i - \frac{1}{2} \sum_{i=1}^m \sum_{j=1}^m \alpha_i \alpha_j y^{(i)} y^{(j)} \langle x^{(i)}, x^{(j)} \rangle$
subject to: $\alpha_i \geq 0, \forall 1 \leq i \leq m; \sum_{i=1}^m \alpha_i y^{(i)} = 0$
- Results of solving Lagrange:
 $\theta = \sum_{i=1}^m \alpha^{(i)} x^{(i)} y^{(i)}$ & $\sum_{i=1}^m \alpha^{(i)} y^{(i)} = 0$
- KTT condition yields:
 $\alpha^{(i)} [y^{(i)}(\theta^T x^{(i)} + \theta_0) - 1] = 0, \forall 1 \leq i \leq m$
- Prediction: $h_{\theta}(x) = \text{sign}(\sum_{i=1}^m \alpha^{(i)} y^{(i)} \langle x^{(i)}, x \rangle + \theta_0)$

3.3.2 Non-separable case

- Optimization problem: $\min_{\theta, \theta_0} \frac{1}{2} \|\theta\|^2 + C \sum_{i=1}^m \xi^{(i)}$, subject to $y^{(i)}(\theta^T x^{(i)} + \theta_0) \geq 1 - \xi^{(i)}, \forall 1 \leq i \leq m$; and $\xi^{(i)} \geq 0, \forall 1 \leq i \leq m$
- Dual problem:
 $\max_{\alpha} \sum_{i=1}^m \alpha_i - \frac{1}{2} \sum_{i=1}^m \sum_{j=1}^m \alpha_i \alpha_j y^{(i)} y^{(j)} \langle x^{(i)}, x^{(j)} \rangle$
subject to: $0 \leq \alpha_i \leq C, \forall 1 \leq i \leq m; \sum_{i=1}^m \alpha_i y^{(i)} = 0$

3.3.3 Hinge loss – another view of SVM

- Define $h_{\theta}(x) = \theta^T x + \theta_0$, rewrite constraints into $\xi^{(i)} = \max(0, 1 - y^{(i)} h_{\theta}(x))$.
- Loss function:
 $J(\theta) = C \sum_{i=1}^m \max(0, 1 - y^{(i)} h_{\theta}(x)) + \frac{1}{2} \|\theta\|^2$

3.3.4 Multiclass SVM

- Loss function:
 $J(\theta) = \sum_{i=1}^m \sum_{j \neq y^{(i)}} \max(0, \theta^{(j)T} x^{(i)} - \theta^{y^{(i)T} x^{(i)} + \Delta)$
- Δ is some fixed margin.

3.3.5 SVM for regression

- Loss function: $L_{\epsilon}(y, \hat{y}) = \max(0, |y - \hat{y}| - \epsilon)$ in
 $J(\theta) = C \sum_{i=1}^m L_{\epsilon}(y^{(i)}, \theta^T x^{(i)} + \theta_0) + \frac{1}{2} \theta^T \theta$
- $J(\theta)$ is convex but not differentiable, the optimization problem is unconstrained.
- Quadratic programming (QP) problem:
 $\min_{\theta} C \sum_{i=1}^m (\xi_+^{(i)} + \xi_-^{(i)}) + \frac{1}{2} \theta^T \theta$, subject to
 $\xi_+^{(i)} \geq 0, \xi_-^{(i)} \geq 0, \forall 1 \leq i \leq m$;
 $\theta^T x^{(i)} + \theta_0 - \xi_-^{(i)} - \epsilon \leq y^{(i)} \leq \theta^T x^{(i)} + \theta_0 + \xi_+^{(i)} + \epsilon$
- Lagrange yields: $\theta = \sum_{i=1}^m (\alpha_+^{(i)} - \alpha_-^{(i)}) x^{(i)}$, where $\alpha_+^{(i)}$ and $\alpha_-^{(i)}$ are Lagrange multipliers of two constraints.
- Prediction: $h_{\theta}(x) = \sum_{i=1}^m (\alpha_+^{(i)} - \alpha_-^{(i)}) \langle x^{(i)}, x \rangle + \theta_0$

4 Neural Networks

Note: A feed forward network with a linear output layer and at least one hidden layer with any squashing activation function (e.g. sigmoid, ReLU), can approximate any function from $\mathbb{R}^d \rightarrow \mathbb{R}$ to arbitrary precision with enough hidden units.

4.1 Activation functions

- $\text{sigmoid}(x) = \sigma(x) = \frac{1}{1+\exp(-x)}$
 $\frac{d}{dx}\sigma(x) = \sigma(x)(1-\sigma(x))$
- $\tanh(x) = \frac{\exp(x)-\exp(-x)}{\exp(x)+\exp(-x)}$ (*unbias, better than sigmoid*)
- $\text{ReLU}(x) = \max(0, x)$
- Softmax function: $g(z_i) = \frac{\exp(z_i)}{\sum_j \exp(z_j)}$ (*output layer*)

4.2 Backpropagation

- Using chain rule, propagate derivatives in inverse order.
- Gradients need to be added up at forks (*accumulative*)!

4.3 Convolutional neural networks

- Let K be filter number, F be filter size, S be stride, P be padding.
- A conv layer takes as input a volume $W_1 \times H_1 \times D_1$, produces an output volume $W_2 \times H_2 \times D_2$:
 $W_2 = \frac{W_1 - F + 2P}{S} + 1$
 $H_2 = \frac{H_1 - F + 2P}{S} + 1$
 $D_2 = K$
- Total number of parameters: $F \times F \times D_1$ weights per filter; $F \times F \times D_1 \times K$ weights, K biases.
- Pooling layer (subsampling): input $W_1 \times H_1 \times D$, output $W_2 \times H_2 \times D$:
 $W_2 = \frac{W_1 - F}{S} + 1$
 $H_2 = \frac{H_1 - F}{S} + 1$
- Classical architecture: [(CONV – RELU)*N – POOL]*M – (FC – RELU)*K – SOFTMAX

4.4 Optimization for training deep models

4.4.1 Momentum

- Require:** learning rate α , initial parameter θ , batch size m , momentum parameter μ , initial velocity v .
- Sample a mini batch of m examples $(x^{(i)}, y^{(i)})$
- Compute gradient estimate $\hat{g} \leftarrow \frac{1}{m} \nabla_{\theta} \sum_i L(h_{\theta}(x^{(i)}), y^{(i)})$
- Compute velocity update $v \leftarrow \mu v - \alpha \hat{g}$
- Apply update $\theta \leftarrow \theta + v$

4.4.2 Nesterov momentum

- Require:** learning rate α , initial parameter θ , batch size m , momentum parameter μ , initial velocity v .
- Sample a mini batch of m examples $(x^{(i)}, y^{(i)})$
- Apply interim update $\tilde{\theta} \leftarrow \tilde{\theta} + \mu v$
- Compute gradient at interim point
 $\hat{g} = \frac{1}{m} \nabla_{\tilde{\theta}} \sum_i L(h_{\tilde{\theta}}(x^{(i)}), y^{(i)})$
- Compute velocity update $v \leftarrow \mu v - \alpha \hat{g}$
- Apply update $\theta \leftarrow \theta + v$

4.4.3 AdaGrad

- Require:** step size α , initial parameter θ , batch size m , $\delta = 10^{-7}$ (constant for numerical stability)
- Compute gradient on minibatch
 $\hat{g}_{ij}^{(l)} \leftarrow \frac{1}{m} \nabla_{\theta_{ij}^{(l)}} \sum_i L(h_{\theta}(x^{(i)}), y^{(i)})$
- Accumulate squared gradient $r_{ij}^{(l)} \leftarrow r_{ij}^{(l)} + \hat{g}_{ij}^{(l)} * \hat{g}_{ij}^{(l)}$
- Compute gradient $\Delta\theta_{ij}^{(l)} \leftarrow -\frac{\alpha}{\delta + \sqrt{r_{ij}^{(l)}}} \hat{g}_{ij}^{(l)}$
- Apply update $\theta_{ij}^{(l)} \leftarrow \theta_{ij}^{(l)} + \Delta\theta_{ij}^{(l)}$

4.4.4 RMSprop

- Require:** step size α , initial parameter θ , batch size m , $\delta = 10^{-7}$ (constant for numerical stability), exponential decay rate ρ .
- Compute gradient on minibatch
 $\hat{g}_{ij}^{(l)} \leftarrow \frac{1}{m} \nabla_{\theta_{ij}^{(l)}} \sum_i L(h_{\theta}(x^{(i)}), y^{(i)})$
- Accumulate squared gradient $r_{ij}^{(l)} \leftarrow \rho r_{ij}^{(l)} + (1 - \rho) \hat{g}_{ij}^{(l)} * \hat{g}_{ij}^{(l)}$
- Compute gradient $\Delta\theta_{ij}^{(l)} \leftarrow -\frac{\alpha}{\delta + \sqrt{r_{ij}^{(l)}}} \hat{g}_{ij}^{(l)}$
- Apply update $\theta_{ij}^{(l)} \leftarrow \theta_{ij}^{(l)} + \Delta\theta_{ij}^{(l)}$

4.4.5 Adam (adaptive moments)

- Require:** step size α (10^{-3} default), initial parameter θ , batch size m , exponential decay rates for moment estimates ρ_1 (0.99 default) & ρ_2 (0.999 default), $\delta = 10^{-7}$ (constant for numerical stability)
- Initialize time step $t = 0$, first and second moment $s = 0$, $r = 0$
- Compute gradient on minibatch
 $\hat{g}_{ij}^{(l)} \leftarrow \frac{1}{m} \nabla_{\theta_{ij}^{(l)}} \sum_i L(h_{\theta}(x^{(i)}), y^{(i)})$
- $t \leftarrow t + 1$
- Update biased first moment $s_{ij}^{(l)} \leftarrow \rho_1 s_{ij}^{(l)} + (1 - \rho_1) \hat{g}_{ij}^{(l)}$
- Update biased second moment
 $r_{ij}^{(l)} \leftarrow \rho_2 r_{ij}^{(l)} + (1 - \rho_2) \hat{g}_{ij}^{(l)} * \hat{g}_{ij}^{(l)}$
- Correct bias in first moment $\hat{s}_{ij}^{(l)} = \frac{s_{ij}^{(l)}}{1 - \rho_1^t}$
- Correct bias in second moment $\hat{r}_{ij}^{(l)} = \frac{r_{ij}^{(l)}}{1 - \rho_2^t}$
- Compute gradient $\Delta\theta_{ij}^{(l)} \leftarrow -\frac{\alpha \hat{s}_{ij}^{(l)}}{\delta + \sqrt{\hat{r}_{ij}^{(l)}}}$
- Apply update $\theta_{ij}^{(l)} \leftarrow \theta_{ij}^{(l)} + \Delta\theta_{ij}^{(l)}$

5 Decision trees

5.1 Cost functions

- Misclassification rate: (\hat{y} = the majority label in \mathcal{D})
 $\text{cost}(\mathcal{D}) = \frac{1}{|\mathcal{D}|} \sum_{(x,y) \in \mathcal{D}} I(y \neq \hat{y})$
- Entropy: (p = fraction of positive examples in \mathcal{D})
 $\text{cost}(\mathcal{D}) = -p \log_2 p - (1 - p) \log_2 (1 - p)$
- Gini index: (same as entropy)
 $\text{cost}(\mathcal{D}) = 2p(1 - p)$

5.2 Decision tree for regression

- Define $\text{cost}(\mathcal{D}) = \sum_{i=1}^m (y^{(i)} - \bar{y})^2$
where $\bar{y} = \frac{1}{|\mathcal{D}|} \sum_{i=1}^m y^{(i)}$

5.3 Node is not worth splitting when

- Node is pure
- Depth exceeds max depth
- $|\mathcal{D}_{\text{left}}|$ or $|\mathcal{D}_{\text{right}}|$ is too small
- Reduction in cost is too small

5.4 Avoid overfitting

- Early stopping – stop growing the tree when the decrease in error is not sufficient to justify the complexity of an additional level.
- Post pruning – grow the full tree and then prune using a validation set to guide subtree removal. Evaluate CV error on each subtree and pick tree whose error is within 1 standard deviation of minimum.

6 Ensembles

6.1 Bagging

- Assume errors of individual members are uncorrelated.
- For regression: when $h_l(x) = f(x) + \epsilon_l(x)$ for $1 \leq l \leq L$, and $\epsilon_l \sim \mathcal{N}(0, \sigma_l^2)$
 $\Rightarrow E_{\text{bag}} = \frac{1}{L} E_{\text{avg}}$ (*expected squared error*)
- For classification: ϵ = error of each classifier, and $\epsilon < \frac{1}{2}$
 $\Rightarrow E_{\text{bag}} = \sum_{i=\frac{L}{2}+1}^L \binom{L}{i} \epsilon^i (1 - \epsilon)^{L-i}$

6.2 Boosting

- Loss function: $J_l = \sum_{i=1}^m w_l^{(i)} I(h_l(x^{(i)}) \neq y^{(i)})$
- Prediction: $h(x) = \text{sign}(\sum_{l=1}^L \alpha_l h_l(x))$
- Adaboost algorithm: initialize $w_1^{(i)} = \frac{1}{m}$, for $1 \leq i \leq m$
 - fit h_l to minimize $J_l = \frac{1}{m} \sum_{i=1}^m w_l^{(i)} L(y^{(i)}, h_l(x^{(i)}))$
 - calculate error rate $\epsilon_l = \frac{\sum_{i=1}^m w_l^{(i)} I(h_l(x^{(i)}) \neq y^{(i)})}{\sum_{i=1}^m w_l^{(i)}}$
 - calculate $\alpha_l = \frac{1}{2} \log\left(\frac{1 - \epsilon_l}{\epsilon_l}\right)$, stop if $\epsilon_l \geq \frac{1}{2}$
 - update $w_{l+1}^{(i)} = \begin{cases} w_l^{(i)} \exp(\alpha_l), & \text{incorrect on } x^{(i)} \\ w_l^{(i)} \exp(-\alpha_l), & \text{correct on } x^{(i)} \end{cases}$

6.3 Gradient boosting

6.3.1 Gradient boosting for regression

- Residuals are negative gradients (squared error loss):
 $J = \sum_{i=1}^m \frac{1}{2} (y - h(x))^2 \Rightarrow \frac{\partial J}{\partial h(x^{(i)})} = h(x^{(i)}) - y^{(i)}$

6.3.1 Gradient boosting for classification

- Loss function: $J = \frac{1}{m} \sum_{i=1}^m D_{KL}(y^{(i)}, h(x^{(i)}))$
- Gradient boosting algorithm (k classes):
 - Start with an initial $h^0 \dots h^k$ for $x^{(1)} \dots x^{(m)}$
 - Repeat until convergence: Calculate matrix of gradients, fit each h_{add} to the negative gradient, $h \leftarrow h + h_{\text{add}}$

7 Probabilistic graphical models

7.1 Directed models – Bayesian network

- $P(X) = \prod_i P(x_i | \text{Parents}(x_i))$
- Reduce number of parameters $O(k^n) \rightarrow O(nk^m)$ if each variable in graph has no more than m parents.

7.2 Undirected models – Markov network

- $\tilde{P}(X) = \prod_{C \in \mathcal{G}} \phi(C)$ (C is a clique in graph)
- Partition function: $Z = \int_X \tilde{P}(X) dX$ or $Z = \sum_X \tilde{P}(X)$
- $P(X) = \frac{1}{Z} \tilde{P}(X)$
- Energy function E : $\tilde{P}(X) = \exp(-E(X))$
high (low) energy \Leftrightarrow low (high) $\tilde{P}(X)$

7.3 Sampling

7.3.1 Ancestral sampling

- For directed graphical models, polynomial time.
- Algorithm:
 - Sort variables in topological order
 - Sample x_i from distribution $P(x_i | \text{Parents}(x_i))$

7.3.2 Gibbs sampling

- For undirected graphical models.
- Algorithm:
 - Start with randomly generated values x_1, \dots, x_n
 - Iteratively visit each x_i and sample a value for it based on $P(x_i | \text{Neighbors}(x_i))$
 - Repeat previous step, generate stream of samples

7.4 Hidden Markov models

- Specified by sets S (hidden states), O (observations) and probability parameters $\lambda = [\pi, a, b]$
 - π is initial state probability
 - a is hidden state transition probability
 - b is emission probability
- Inference problems:
 - Filtering: $P(X_t | e_1, \dots, e_t)$
 - Smoothing: $P(X_k | e_1, \dots, e_t), k < t$
 - Most likely state sequence:
 $\arg \max_{x_1, \dots, x_t} P(X_1, \dots, X_t | e_1, \dots, e_t)$

7.4.1 Forward computation – filtering

- Define: $\alpha_t(i) = P(e_1, \dots, e_t, X_t = s_i)$
- Algorithm:
 - $\alpha_0(i) = \pi_i, 1 \leq i \leq n$ where $|S| = n$
 - $\alpha_{t+1}(i) = b_j(e_{t+1}) \sum_{j=1}^n \alpha_t(j) a_{ij}, 1 \leq j \leq n, 0 \leq t \leq T-1$
- Time complexity: $O(n^2T)$

7.4.2 Backward computation – smoothing

- $P(X_k | e_1, \dots, e_t) \propto P(e_{k+1}, \dots, e_t | X_k) P(X_k | e_1, \dots, e_t)$ where $t > k$
- Define: $\beta_k(i) = P(e_{k+1}, \dots, e_t | X_k = s_i)$
- Algorithm:
 - $\beta_T(i) = 1, 1 \leq i \leq n$
 - $\beta_k(i) = \sum_{j=1}^n a_{ij} b_j(e_{k+1}) \beta_{k+1}(j), 1 \leq j \leq n, 0 \leq k \leq T-1$
- Time complexity: $O(n^2T)$

7.4.3 Viterbi algorithm – most likely sequence

- Define:
 $\delta_t(i) = \max_{x_1, \dots, x_{t-1}} P(X_1, \dots, X_{t-1}, X_t = s_i, e_1, \dots, e_t)$
- Algorithm:
 - $\delta_0(i) = \pi(i), 1 \leq i \leq n$
 - $\delta_{t+1}(j) = \max_i \delta_t(i) a_{ij} b_j(e_{t+1}), 1 \leq j \leq n, 0 \leq t \leq T-1$
- Time complexity: $O(n^2T)$

7.4.4 Parameter estimation

- Paired sequences: $\hat{a}_{ij} = \frac{\#s_i \rightarrow s_j}{\#s_i}$ & $\hat{b}_j(e_k) = \frac{\#s_j \rightarrow e_k}{\#s_j}$
- Observation sequences only – Baum-Welch EM:
 - Define: $\xi_t(i, j) = P(X_t = s_i, X_{t+1} = s_j | e_1, \dots, e_T, \lambda)$
 $\xi_t(i, j) = \frac{\alpha_t(i) a_{ij} b_j(e_{t+1}) \beta_{t+1}(j)}{\sum_{i=1}^n \sum_{j=1}^n \alpha_t(i) a_{ij} b_j(e_{t+1}) \beta_{t+1}(j)}$
 - Let $\gamma_t(i) = P(X_t = s_i | e_1, \dots, e_T, \lambda) = \sum_{j=1}^n \xi_t(i, j)$
 - Estimate $\hat{\pi}_i = \gamma_1(i)$, $\hat{a}_{ij} = \frac{\sum_{t=1}^{T-1} \xi_t(i, j)}{\sum_{t=1}^{T-1} \gamma_t(i)}$,
 $\hat{b}_j(e_k) = \frac{\sum_{t=1}^T \gamma_t(j) * I(E_t = e_k)}{\sum_{t=1}^T \gamma_t(j)}$
 - Algorithm:
 - Guess $\lambda_0 = [\pi_0, a_0, b_0]$
 - Repeat until convergence:
 - Calculate α, β from λ
 - Re-estimate λ from α, β

8 Unsupervised learning

8.1 Principal components analysis (PCA)

- Assume: data distribution is unimodal Gaussian (fully explained by mean & variance)
- Assume: information to be preserved is in the variance
- Project data $\mathbb{R}^d \rightarrow \mathbb{R}^k$ ($k < d$), maximizing variance.

8.1.1 PCA method

- Zero-mean, unit variance transform on \mathcal{D}
- Find S = covariance matrix of transformed \mathcal{D}
- Find $\lambda_1, \dots, \lambda_k$ (the k largest eigenvalues of S) and associated eigenvectors u_1, \dots, u_k
- Project $x^{(i)} \mapsto [u_1^T x^{(i)}, \dots, u_k^T x^{(i)}]^T$, where $x^{(i)} \in \mathbb{R}^d$

8.1.2 Kernel PCA

- Idea: map $x^{(i)} \mapsto \phi(x^{(i)})$ where $\phi: \mathbb{R}^d \rightarrow \mathbb{R}^D$ ($D \gg d$)
- The result when project back to \mathbb{R}^d will be nonlinear.
- Algorithm:
 - Pick a kernel
 - Construct kernel matrix K over data $x^{(1)}, \dots, x^{(m)}$
 - Zero mean the kernel matrix K to get \tilde{K}
 - Solve the eigenvalue problem $\tilde{K}\alpha = \lambda\alpha, \alpha \in \mathbb{R}^m$
 - For a new x , we project it as: $y_j = \sum_{i=1}^m \alpha_j^{(i)} \kappa(x, x^{(i)})$, for $j = 1, \dots, L$ (# of components) where eigenvectors are ordered by value

8.2 Expectation maximization algorithm (EM)

8.2.1 K-means

- Cost function: $J = \sum_{i=1}^m \sum_{k=1}^K z_k^{(i)} \|x^{(i)} - \mu_k\|^2$
- E-step: cluster assignment, minimize J wrt. z , fix μ
- M-step: relocate means, minimize J wrt. μ , fix z
- Time complexity: $O(mK)$ per iteration
- Converges to local minimum & vulnerable to outliers.

8.2.2 Gaussian mixture model (GMM)

- Generative mode:
 $P(x^{(i)}) = \sum_{k=1}^K P(z^{(i)} = k) P(x^{(i)} | z^{(i)} = k)$
 $z^{(i)} \sim \text{Multinomial}(\pi); \pi_k > 0, \sum_k \pi_k = 1$
 $x^{(i)} |_{z^{(i)}=k} \sim \mathcal{N}(\mu_k, \Sigma_k)$
- Infer $z^{(i)}$ for each $x^{(i)}$, where $\theta = \{\pi, \mu, \Sigma\}$:
 $P(z^{(i)} = k | x^{(i)}; \theta) = \frac{P(z^{(i)}=k) P(x^{(i)} | z^{(i)}=k; \theta)}{\sum_{k'} P(z^{(i)}=k') P(x^{(i)} | z^{(i)}=k'; \theta)}$
- Soft EM algorithm:
 - Guess values of $\theta = \{\pi, \mu, \Sigma\}$
 - E-step: calculate the responsibility of each component toward generating $x^{(i)}$: $r_k^{(i)} = P(z^{(i)} = k | x^{(i)}; \theta)$
 - M-step: given $r_k^{(i)}$ and $x^{(i)}, 1 \leq i \leq m, 1 \leq k \leq K$, re-estimate π, μ, Σ :
 $\pi_k = \frac{1}{m} \sum_{i=1}^m r_k^{(i)}, \quad \mu_k = \frac{\sum_{i=1}^m r_k^{(i)} x^{(i)}}{\sum_{i=1}^m r_k^{(i)}}$
 $\Sigma_k = \frac{\sum_{i=1}^m r_k^{(i)} (x^{(i)} - \mu_k)(x^{(i)} - \mu_k)^T}{\sum_{i=1}^m r_k^{(i)}}$

9 Reinforcement learning

9.1 Markov decision process (MDP)

9.1.1 The model

- A set of states S , a subset of which are terminal states.
- Actions(s): possible actions from state s , no actions from terminal states.
- A transition function $T(s, a, s')$: probability of transitioning to state s' if action a is taken in state s .
- A reward function $r(s, a, s')$: reward for taking action a in state s and ending up in state s' , or $r(s, a)$ or $r(s)$.

9.1.2 The value function – expected utility

- Definition: $V_\pi(s) = E[\sum_{t=0}^{\infty} r(s_t, \pi(s_t)) | \pi, s_0 = s]$
- Recursive:
 $V_\pi(s) = \sum_{s'} T(s, \pi(s), s') [r(s, \pi(s), s') + V_\pi(s')]$
- Q-function: $Q_\pi(s, a) = \sum_{s'} T(s, a, s') [r(s, a, s') + V_\pi(s')]$

9.1.3 Policy evaluation

- Approach 1: set up linear equations and solve for V_π
- Approach 2: (iterative improvement) $O(|S|)$ each iteration
 - Repeat until convergence:
 - For each state s :
 $V_\pi^i(s) \leftarrow \sum_{s'} T(s, \pi(s), s') [r(s, \pi(s), s') + V_\pi^{i-1}(s')]$

9.1.4 Optimality

- Definition: $V^*(s) = V_{\pi^*}(s) = \max_\pi V_\pi(s)$
- Bellman's equation:
 $V^*(s) = \max_{a \in \text{Actions}(s)} T(s, a, s') [r(s, a, s') + V^*(s')]$
- Given V^* : $Q^*(s, a) = \sum_{s'} T(s, a, s') [r(s, a, s') + V^*(s)]$
- Given Q^* : $V^*(s) = \max_{a \in \text{Action}(s)} Q^*(s, a)$

9.2 Solving MDPs

9.2.1 Policy iteration

- Policy improvement:
 - Compute $Q_\pi(s, a)$ from $V_\pi(s)$
 - Update π : $\pi'(s) = \arg \max_{a \in \text{Action}(s)} Q(s, a)$
- Policy iteration algorithm:
 - Start with a random policy π
 - Repeat until no change to policy occurs:
 - Compute value of policy π (policy evaluation)
 - Improve the policy at each state (policy improvement)

9.2.2 Value iteration

- Note: no explicit policy.
- Value iteration algorithm:
 - Start with $V^{(0)}(s) = 0$ for all states s in S
 - Repeat until convergence:
 - Bellman update: $V^i(s) \leftarrow \max_{a \in \text{Action}(s)} \sum_{s'} T(s, a, s') [r(s, a, s') + V^{i-1}(s')]$

9.3 Model-based RL

9.4 Model-free RL

9.5 TD learning

9.6 Q learning

Appendix

Distributions

- Poisson distribution PMF: $P(X = k) = \frac{\lambda^k e^{-\lambda}}{k!}$
- Normal distribution PDF:

$$f(x; \mu, \sigma) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp \left\{ -\frac{(x-\mu)^2}{2\sigma^2} \right\}$$
- Multivariate normal distribution PDF:

$$f(X; \mu, \Sigma) = \frac{1}{\sqrt{(2\pi)^k |\Sigma|}} \exp \left(-\frac{1}{2} (X - \mu)^T \Sigma^{-1} (X - \mu) \right)$$
- Laplace distribution PDF: $f(x; \mu, b) = \frac{1}{2b} \exp \left\{ -\frac{|x-\mu|}{b} \right\}$
- Beta distribution PDF: $f(x; \alpha, \beta) = \frac{1}{B(\alpha, \beta)} x^{\alpha-1} (1-x)^{\beta-1}$
- Given that $X_1 \sim \text{Pois}(\lambda_1)$, $X_2 \sim \text{Pois}(\lambda_2)$, $X = X_1 + X_2$, X_1 and X_2 are independent: $X \sim \text{Pois}(\lambda_1 + \lambda_2)$.
- Given that $P(X_0 = x_0) = \alpha_0 \exp \left\{ -\frac{(x_0 - \mu_0)^2}{2\sigma_0^2} \right\}$, and that

$$P(X_1 = x_1 | X_0 = x_0) = \alpha_1 \exp \left\{ -\frac{(x_1 - x_0)^2}{2\sigma_1^2} \right\};$$

$$P(X_1 = x_1) = \alpha_0 \alpha_1 \sqrt{\frac{2\pi\sigma_0^2\sigma_1^2}{\sigma_0^2 + \sigma_1^2}} \exp \left\{ -\frac{1}{2} \frac{(x_1 - \mu_0)^2}{\sigma_0^2 + \sigma_1^2} \right\}$$

Information theory

- Conditional information:

$$H(Y|X) = \sum_{y \in Y} \sum_{x \in X} p(x, y) \log \frac{p(x)}{p(x, y)}$$
- Mutual information:

$$I(X; Y) = \sum_{y \in Y} \sum_{x \in X} p(x, y) \log \frac{p(x, y)}{p(x)P(y)}$$
- Kullback-Leibler (KL) divergence:

$$D_{KL}(p||q) = \sum_x p(x) \log \frac{p(x)}{q(x)}$$
- Cross entropy: $H(p, q) = E_p[-\log q] = H(p) + D_{KL}(p||q)$

Convex function

- A function $f(x)$ is convex on a set S iff for $\lambda \in [0, 1]$, and $\forall x, y \in S$: $f(\lambda x + (1 - \lambda)y) \leq \lambda f(x) + (1 - \lambda)f(y)$.
- A function $f(x)$ is convex on a set S iff $\frac{d^2}{dx^2} f(x)$ is positive semidefinite everywhere in the set.

Convergence of perceptron

Proof. Let $\theta^{(k-1)}$ be the parameter vector when the algorithm makes a mistake on (x, y) .

$$\theta^{(k)} = \theta^{(k-1)} + \eta xy$$

Take dot product on both sides with θ^* (some separating hyperplane)

$$\begin{aligned} \theta^{*T} \theta^{(k)} &= \theta^{*T} (\theta^{(k-1)} + \eta xy) \\ &= \theta^{*T} \theta^{(k-1)} + \eta y (\theta^{*T} x) \\ &\geq \theta^{*T} \theta^{(k-1)} + \eta \gamma \end{aligned}$$

If $\theta^{(0)}$ is all zeros vector, then

$$\theta^{*T} \theta^{(k)} \geq \eta k \gamma \quad (1)$$

From the update rule,

$$\begin{aligned} \theta^{(k)} &= \theta^{(k-1)} + \eta xy \\ \|\theta^{(k)}\|^2 &= \|\theta^{(k-1)} + \eta xy\|^2 \\ &= \|\theta^{(k-1)}\|^2 + \eta^2 y^2 \|x\|^2 + 2\eta y (\theta^{(k-1)T} x) \\ &\leq \|\theta^{(k-1)}\|^2 + \eta^2 \|x\|^2 \\ &\leq \|\theta^{(k-1)}\|^2 + \eta^2 R^2 \end{aligned}$$

Starting with $\theta^{(0)}$ of all zeros,

$$\|\theta^{(k-1)}\|^2 \leq k \eta^2 R^2 \quad (2)$$

Putting (1) and (2) together, $k \leq \frac{R^2 \|\theta^*\|^2}{\gamma^2}$. \square

Mercer's theorem

Proof. Since K is positive definite, $K = u^T \Lambda u$, where Λ is a diagonal matrix with entries $\lambda^{(i)} > 0$. Consider an element $\kappa(x^{(i)}, x^{(j)})$ of K . We can construct this element as follows

$$\kappa(x^{(i)}, x^{(j)}) = (\Lambda^{\frac{1}{2}} u_{:,i})^T (\Lambda^{\frac{1}{2}} u_{:,j})$$

Now define $\phi(x^{(i)}) = \Lambda^{\frac{1}{2}} u_{:,i}$, then we have

$$\kappa(x^{(i)}, x^{(j)}) = \phi(x^{(i)})^T \phi(x^{(j)}).$$

Adaboost principle

Kernel PCA derivation

Correctness of EM

Proof. For any \mathbf{Z} with non-zero probability $p(\mathbf{Z}|\mathbf{X}, \theta)$, we can write:

$$\log p(\mathbf{X}|\theta) = \log p(\mathbf{X}, \mathbf{Z}|\theta) - \log p(\mathbf{Z}|\mathbf{X}, \theta)$$

We take the expectation over possible values of the unknown data \mathbf{Z} under the current parameter estimate $\theta^{(t)}$ by multiplying both sides by $p(\mathbf{Z}|\mathbf{X}, \theta^{(t)})$ and summing over \mathbf{Z} . The left-hand side is the expectation of a constant, so we get:

$$\begin{aligned} \log p(\mathbf{X}|\theta) &= \sum_{\mathbf{Z}} p(\mathbf{Z}|\mathbf{X}, \theta^{(t)}) \log p(\mathbf{X}, \mathbf{Z}|\theta) \\ &\quad - \sum_{\mathbf{Z}} p(\mathbf{Z}|\mathbf{X}, \theta^{(t)}) \log p(\mathbf{Z}|\mathbf{X}, \theta) \\ &= Q(\theta|\theta^{(t)}) + H(\theta|\theta^{(t)}) \end{aligned}$$

where $H(\theta|\theta^{(t)})$ is defined by the negated sum it is replacing. This last equation holds for any value of θ including $\theta = \theta^{(t)}$,

$$\log p(\mathbf{X}|\theta^{(t)}) = Q(\theta^{(t)}|\theta^{(t)}) + H(\theta^{(t)}|\theta^{(t)})$$

and subtracting this last equation from the previous equation gives

$$\begin{aligned} \log p(\mathbf{X}|\theta) - \log p(\mathbf{X}|\theta^{(t)}) &= Q(\theta|\theta^{(t)}) - Q(\theta^{(t)}|\theta^{(t)}) \\ &\quad + H(\theta|\theta^{(t)}) - H(\theta^{(t)}|\theta^{(t)}) \end{aligned}$$

However, Gibbs' inequality tells us that $H(\theta|\theta^{(t)}) \geq H(\theta^{(t)}|\theta^{(t)})$, so we can conclude that

$$\log p(\mathbf{X}|\theta) - \log p(\mathbf{X}|\theta^{(t)}) \geq Q(\theta|\theta^{(t)}) - Q(\theta^{(t)}|\theta^{(t)})$$

In words, choosing θ to improve $Q(\theta|\theta^{(t)})$ beyond $Q(\theta^{(t)}|\theta^{(t)})$ cannot cause $\log p(\mathbf{X}|\theta)$ to decrease below $\log p(\mathbf{X}|\theta^{(t)})$, and so the marginal likelihood of the data is non-decreasing. \square